

cnn_model_1-Fixed_for_final

February 16, 2020

```
In [8]: import matplotlib.pyplot as plt
        from keras import applications
        from keras.preprocessing.image import ImageDataGenerator
        from keras import optimizers
        from keras.models import Sequential
        from keras.layers import Dropout, Flatten, Dense, Conv2D, MaxPooling2D
        from keras.applications.inception_v3 import InceptionV3
        from keras.preprocessing import image
        from keras.models import Model
        from keras.layers import Dense, Flatten
        from keras import backend as K
        import numpy as np
        import pandas as pd
        import os
        from sklearn.metrics import classification_report, confusion_matrix
        import sklearn.metrics as metrics
        import sklearn
        from sklearn.metrics import roc_auc_score
        from sklearn.metrics import roc_curve
        import matplotlib.pyplot as plt
        %matplotlib inline

In [9]: train_data_dir = '/home/mlab/Documents/brats_hl_data/train'
        validation_data_dir = '/home/mlab/Documents/brats_hl_data/val'
        # 44938
        # 5616
        nb_train_samples = 44938
        nb_validation_samples = 5616
        epochs = 8
        batch_size = 128
        # prepare data augmentation configuration
        train_datagen = ImageDataGenerator(
            rescale=1. / 255,
            shear_range=0.2,
            zoom_range=0.2,
            horizontal_flip=True)

        test_datagen = ImageDataGenerator(rescale=1. / 255)
```

```

train_generator = train_datagen.flow_from_directory(
    train_data_dir,
    target_size=(150, 150),
    batch_size=batch_size,
    class_mode='binary')

validation_generator = test_datagen.flow_from_directory(
    validation_data_dir,
    target_size=(150, 150),
    batch_size=batch_size,
    class_mode='binary')

```

Found 44938 images belonging to 2 classes.
Found 5616 images belonging to 2 classes.

```

In [10]: #uses a smal vggnet
         #32->64->128->512
         model = Sequential()
         model.add(Conv2D(32,(3,3),activation='relu',input_shape=(150,150,3)))
         model.add(MaxPooling2D((2,2)))
         model.add(Conv2D(64,(3,3),activation='relu'))
         model.add(MaxPooling2D((2,2)))
         model.add(Conv2D(128,(3,3),activation='relu'))
         model.add(MaxPooling2D((2,2)))
         model.add(Conv2D(512,(3,3),activation='relu'))
         model.add(MaxPooling2D((2,2)))
         model.add(Flatten())
         model.add(Dropout(0.5))
         model.add(Dense(512,activation='relu'))
         model.add(Dense(1,activation='sigmoid')) #uses sigmoid at the end because we onlu have two classes

```

W0216 09:35:28.779163 140375213623104 deprecation_wrapper.py:119] From /home/mlab/anaconda3/lib/python

W0216 09:35:28.817338 140375213623104 deprecation_wrapper.py:119] From /home/mlab/anaconda3/lib/python

W0216 09:35:28.822030 140375213623104 deprecation.py:506] From /home/mlab/anaconda3/lib/python3.7/site-p

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

```

In [11]: model.summary()

```

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 148, 148, 32)	896

max_pooling2d_1 (MaxPooling2 (None, 74, 74, 32) 0

conv2d_3 (Conv2D) (None, 72, 72, 64) 18496

max_pooling2d_2 (MaxPooling2 (None, 36, 36, 64) 0

conv2d_4 (Conv2D) (None, 34, 34, 128) 73856

max_pooling2d_3 (MaxPooling2 (None, 17, 17, 128) 0

conv2d_5 (Conv2D) (None, 15, 15, 512) 590336

max_pooling2d_4 (MaxPooling2 (None, 7, 7, 512) 0

flatten_1 (Flatten) (None, 25088) 0

dropout_1 (Dropout) (None, 25088) 0

dense_1 (Dense) (None, 512) 12845568

dense_2 (Dense) (None, 1) 513

=====

Total params: 13,529,665

Trainable params: 13,529,665

Non-trainable params: 0

In [12]: #We use the RMSprop optimizer with a learning rate of 0.0001

#We use binary_crossentropy loss because its a binary classification

#optimizer = Adam(learning_rate=lr_schedule(0))

model.compile(loss='binary_crossentropy',optimizer=optimizers.Adam(lr=1e-4),metrics=['acc'])

W0216 09:35:28.871100 140375213623104 deprecation_wrapper.py:119] From /home/mlab/anaconda3/lib/python

W0216 09:35:28.874930 140375213623104 deprecation_wrapper.py:119] From /home/mlab/anaconda3/lib/python

W0216 09:35:28.878212 140375213623104 deprecation.py:323] From /home/mlab/anaconda3/lib/python3.7/site-p

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

In [13]: true_classes = train_generator.classes

print(true_classes)

class_labels = list(train_generator.class_indices.keys())

print(class_labels)

[0 0 0 ... 1 1 1]

['high', 'low']

```

In [14]: # train the model on the new data for a few epochs
         history =model.fit_generator(train_generator,
                                     steps_per_epoch=nb_train_samples//batch_size,
                                     epochs=epochs,
                                     validation_data=validation_generator,
                                     validation_steps=nb_validation_samples//batch_size)

Epoch 1/8
351/351 [=====] - 568s 2s/step - loss: 0.6214 - acc: 0.6449 - val_loss: 0.5550
Epoch 2/8
351/351 [=====] - 559s 2s/step - loss: 0.5550 - acc: 0.7023 - val_loss: 0.5288
Epoch 3/8
351/351 [=====] - 557s 2s/step - loss: 0.5288 - acc: 0.7244 - val_loss: 0.5031
Epoch 4/8
351/351 [=====] - 558s 2s/step - loss: 0.5031 - acc: 0.7394 - val_loss: 0.4801
Epoch 5/8
351/351 [=====] - 558s 2s/step - loss: 0.4801 - acc: 0.7558 - val_loss: 0.4581
Epoch 6/8
351/351 [=====] - 558s 2s/step - loss: 0.4581 - acc: 0.7704 - val_loss: 0.4400
Epoch 7/8
351/351 [=====] - 557s 2s/step - loss: 0.4400 - acc: 0.7801 - val_loss: 0.4195
Epoch 8/8
351/351 [=====] - 557s 2s/step - loss: 0.4195 - acc: 0.7948 - val_loss: 0.4195

In [15]: # # #save the model
         # model.save_weights('model_weights.h5')
         # model.save('model_keras.h5')

In [16]: #Confution Matrix and Classification Report
         Y_pred = model.predict_generator(validation_generator, nb_validation_samples // batch_size+1)

In [17]: # y_pred = np.argmax(Y_pred, axis=1)
         y_pred = (Y_pred<0.475).astype(np.int)

         # print('Confusion Matrix')
         # print(confusion_matrix(true_classes_1, y_pred))
         # print('Classification Report')
         # print(classification_report(validation_generator.classes, y_pred,
         #                             target_names=class_labels_1))

In [19]: true_classes_1 = validation_generator.classes
         print(true_classes)
         class_labels_1 = list(validation_generator.class_indices.keys())
         print(class_labels_1)

[0 0 0 ... 1 1 1]
['high', 'low']

```

```
In [20]: confusion_matrix = metrics.confusion_matrix(true_classes_1,y_pred)
        print(confusion_matrix)
```

```
[[1563 1212]
 [1618 1223]]
```


```
In [21]: report= sklearn.metrics.classification_report(true_classes_1, y_pred,
        target_names = class_labels_1)

        print(report)
```

	precision	recall	f1-score	support
high	0.49	0.56	0.52	2775
low	0.50	0.43	0.46	2841
micro avg	0.50	0.50	0.50	5616
macro avg	0.50	0.50	0.49	5616
weighted avg	0.50	0.50	0.49	5616

```
In [22]: fpr, tpr, thresholds = roc_curve(validation_generator.classes, y_pred)
```

```
# create plot
plt.plot(fpr, tpr, label='ROC curve')
plt.plot([0, 1], [0, 1], 'k--', label='Random guess')
_ = plt.xlabel('False Positive Rate')
_ = plt.ylabel('True Positive Rate')
_ = plt.title('ROC Curve')
_ = plt.xlim([-0.02, 1])
_ = plt.ylim([0, 1.02])
_ = plt.legend(loc="lower right")
```



```
In [23]: roc_auc_score(validation_generator.classes, y_pred)
```

```
Out[23]: 0.49686273390602853
```

```
In [24]: batchX, batchy = train_generator.next()
_, accuracy = model.evaluate(batchX, batchy)
print('Accuracy training: %.2f' % (accuracy*100))
batchXv, batchyv = validation_generator.next()
_, accuracy = model.evaluate(batchXv, batchyv)
print('Accuracy val: %.2f' % (accuracy*100))
```


```
128/128 [=====] - 0s 3ms/step
Accuracy training: 80.47
128/128 [=====] - 0s 3ms/step
Accuracy val: 78.91
```

```
In [25]: #plot the train and val curve
#get the details from the history object
acc = history.history['acc']
val_acc=history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
```


```
epochs = range(1,len(acc)+1)

#train and validation accuracy
plt.plot(epochs,acc,'b',label='Training accuracy')
plt.plot(epochs,val_acc,'r',label='Validation accuracy')
plt.title('Training and Validation accuracy')
plt.legend()
```

Out[25]: <matplotlib.legend.Legend at 0x7fab34e13da0>



```
In [26]: #train and validation loss
plt.plot(epochs, loss, 'b',label='Training loss')
plt.plot(epochs, val_loss, 'r',label='Validation loss')
plt.title('Training and Validation loss')
plt.legend()
plt.show()
```



```
In [27]: test_generator = test_datagen.flow_from_directory('/home/mlab/Documents/brats_hl_data/test',
class_mode='binary',
batch_size=batch_size,
target_size=(150,150))
```

Found 5619 images belonging to 2 classes.

```
In [28]: scores = model.evaluate_generator(test_generator, steps=nb_validation_samples//batch_size)
```

```
In [29]: print("%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
```

acc: 81.16%

```
In [30]: #Confution Matrix and Classification Report
# Y_pred = model.predict_generator(val_generator, 5616 // batch_size)
nb_test_samples=5619
Y_pred1 = model.predict_generator(test_generator,nb_test_samples//batch_size+1)
# y_pred = np.argmax(Y_pred,axis=1)
```

```
In [31]: true_classes_2 = test_generator.classes
print(true_classes_2)
class_labels_2 = list(test_generator.class_indices.keys())
print(class_labels_2)
```



```
[0 0 0 ... 1 1 1]
['high', 'low']
```

```
In [32]: # y_pred1 = (Y_pred1<0.5).astype(np.int)
        y_pred1 = (Y_pred1 < 0.475).astype(np.int)
        # print(y_pred)
        # print('Confusion Matrix')
        # print(confusion_matrix(true_classes_2, y_pred1))
        # print('Classification Report')
        # print(classification_report(true_classes_2, y_pred1, target_names=class_labels_2))
```

```
In [33]: confusion_matrix1 = metrics.confusion_matrix(true_classes_2,y_pred1)
        print(confusion_matrix1)
```

```
[[1574 1202]
 [1660 1183]]
```

```
In [34]: report1= sklearn.metrics.classification_report(true_classes_2, y_pred1,
        target_names = class_labels_2)
        print(report1)
```

	precision	recall	f1-score	support
high	0.49	0.57	0.52	2776
low	0.50	0.42	0.45	2843
micro avg	0.49	0.49	0.49	5619
macro avg	0.49	0.49	0.49	5619
weighted avg	0.49	0.49	0.49	5619

```
In [35]: print(" Loss: ", scores[0],"\n","Accuracy: ", scores[1])
```

```
Loss: 0.39541898078696675
Accuracy: 0.8115915697674418
```